

CS542000 Cloud Programming

Lab2: Hbase, Hive & Pig

Josh Kao

National Tsing Hua University

2015/06/01

Objective

- To get familiar with
 - 1 Using HBase with API
 - 2 Querying data by Hive, Pig

Outline I

- 1 Overview
- 2 HBase
- 3 Hive
- 4 Pig
- 5 Lab Exercise
- 6 Programming Guide
- 7 Reference

What is Hbase?

*“HBase is an open source, **non-relational, distributed database** modeled after Google’s BigTable and written in Java. It is developed as part of Apache Software Foundation’s Apache Hadoop project and **runs on top of HDFS** (Hadoop Distributed Filesystem), providing BigTable-like capabilities for Hadoop.”*

— Wikipedia

What is Hive?


“Apache Hive is a data warehouse infrastructure built on top of Hadoop for providing data summarization, query, and analysis.”
— Wikipedia

What is Pig?

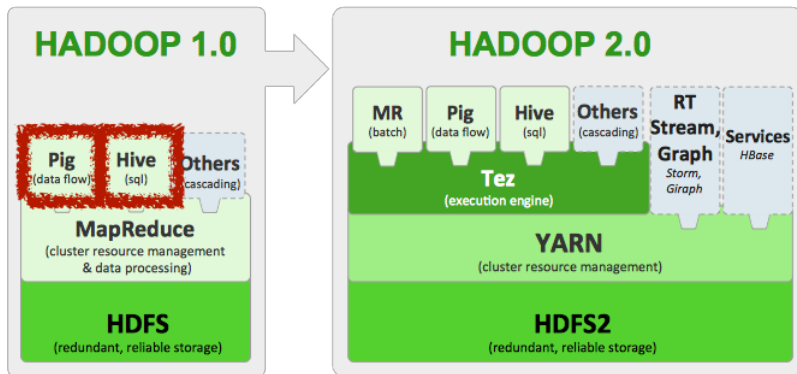
*"Pig is a high-level platform for **creating MapReduce programs** used with Hadoop. The language for this platform is called Pig Latin. Pig Latin abstracts the programming from the Java MapReduce idiom into a notation which makes MapReduce programming high level, similar to that of SQL for RDBMS systems. Pig Latin can be extended **using UDF** (User Defined Functions)."*

— Wikipedia

Hadoop ecosystem evolution



Today's mission




Outline I

- 1 Overview
- 2 HBase**
- 3 Hive
- 4 Pig
- 5 Lab Exercise
- 6 Programming Guide
- 7 Reference

Logical Data Model

- Tables are sorted by **row key**
- Table schema only defines its column families
 - Each family consists of any number of columns
 - Value of each column consists of any number of version (timestamp)
 - Columns can be added dynamically
 - #Columns of each row can be different with other rows
 - Only one data type: `byte[]`

Logical Data Model



Physical Data Model

- Columns in the **same column family** would be group in the **same physical storage**, and would be sorted by column

How to access data?

- Doesn't provide SQL Language
- Access data by
 - ① **getRow()** : get one row range data, besides, user can get specified timestamp data
 - ② **scan()**: get data from whole table

Using HBase

- `$ hbase shell [YOUR_SCRIPT]`
- Cause we share the same database, please **name your table with your ID as the prefix**. e.g. 103065566_helloTable

Scenario

- **List** all tables in HBase

- `> list`

- **Create** a new table

- `> create '{TABLE_NAME}', '{COLUMN_FAMILY1}', ...`

- **Put** value

- `> put '{TABLE_NAME}', '{ROW}',
'{COLUMN_FAMILY[:QUALIFIER]]', '{VALUE}'`

Scenario

- **Get** value

- `> get '{TABLE_NAME}', '{ROW}'[, '{COLUMN1}', ...]`

- **Scan** table

- `> scan '{TABLE_NAME}'[, optional setting]`

Scenario

- **Delete** value

- ```
> delete '{TABLE_NAME}', '{ROW}',
 '{COLUMN_FAMILY[:QUALIFIER]}'
```

- **Remove** table

- ```
> disable '{TABLE_NAME}'
```
- ```
> drop '{TABLE_NAME}'
```

- Learn more from help :D

- ```
> help
```

Why we need Hive/Pig?

- Need high-level languages
 - Writing java programs for everything is verbose and slow
 - Not everyone wants to (or can) write java

Outline I

- 1 Overview
- 2 HBase
- 3 Hive**
- 4 Pig
- 5 Lab Exercise
- 6 Programming Guide
- 7 Reference

Hive

- Query language is **Hive QL**, variant of SQL
- Tables stored on HDFS as flat files
- Not designed for online transaction processing
- Does not offer real-time queries and row level updates.

Primitive Types

- Integers
 - **TINYINT** - 1 byte integer
 - **SMALLINT** - 2 byte integer
 - **INT** - 4 byte integer
 - **BIGINT** - 8 byte integer
- Boolean type
 - **BOOLEAN** - TRUE/FALSE
- Floating point numbers
 - **FLOAT** - single precision
 - **DOUBLE** - double precision
- String type
 - **STRING** - sequence of characters in a specified character set

Complex Types

- Structs
- Maps
- Arrays

Using Hive

- ```
$ hive -hiveconf hbase.master={HBase_MASTER}:{PORT}
--auxpath /opt/hive/lib/hive-hbase-handler-0.11.0.jar,
/opt/hive/lib/hbase-0.94.27.jar,
/opt/hive/lib/zookeeper-3.4.6.jar
[-f {YOUR_SCRIPT}.q]
```
- Learn basic usage of SQL by yourself
  - W3C School
  - Hive Tutorial  
<https://cwiki.apache.org/confluence/display/Hive/Tutorial>

# Load HBase tables to Hive

- *hive* > CREATE EXTERNAL TABLE {TABLE\_NAME}  
> ({SCHEMA}) STORE BY  
> 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
> WITH SERDEPROPERTIES("hbase.columns.mapping" =  
> "{CF1}:{Q1},{CF1}:{Q2},{CF2}: ..."  
> TBLPROPERTIES("hbase.table.name" = "{TABLE\_NAME}")



# Outline I

- 1 Overview
- 2 HBase
- 3 Hive
- 4 Pig**
- 5 Lab Exercise
- 6 Programming Guide
- 7 Reference

# Using Pig

- `$ pig [-x local {YOUR_SCRIPT}.pig]`  
`-Dpig.additional.jars=/opt/hbase/lib/protobuf-java-2.4.0a.jar`
- Pig Tutorial <https://pig.apache.org/docs/r0.7.0/tutorial.html>

# Load HBase tables to Pig

- ```
grunt > {VAR} = LOAD '{HBASE_PATH}'  
    USING org.apache.pig.backend.hbase.HBaseStorage  
    ('{CF1}:{Q1} {CF2}:{Q2} ...', '-loadKey true')  
    AS ({FIELD1}:{TYPE},...);
```
- ```
grunt > math = LOAD 'hbase://103065566_math'
 USING org.apache.pig.backend.hbase.HBaseStorage
 ('grade:name grade:math', '-loadKey true')
 AS (id:CHARARRAY, name:CHARARRAY, math:DOUBLE);
```

# Pig Example

```
> visits = LOAD '/data/visits' AS (user, url, time);
```

```
> visits = FOREACH visits GENERATE user, Canonicalize(url),
time;
```

```
> pages = LOAD '/data/pages' AS (url, pagerank);
```

```
> vp = JOIN visits BY url, pages BY url;
```

```
> userVisits = GROUP vp BY user;
```

## Pig Example (Continue)

```
> userPageranks = FOREACH userVisits GENERATE user,
 AVG(vp.pagerank) AS avgpr;
```

```
> goodUsers = FILTER userPageranks BY avgpr > '0.5';
```

```
> DUMP goodUsers;
```

```
> STORE goodUsers INTO '/data/good_users';
```

# Outline I

- 1 Overview
- 2 HBase
- 3 Hive
- 4 Pig
- 5 Lab Exercise**
- 6 Programming Guide
- 7 Reference

# Problem Description

- Part 1 - Load data to HBase by API
  - ➊ Modify sample code to load file math to HBase
  - ➋ This table named {STUDENT\_ID}\_math with column-family **grade** and qualifier **name,math**
  - ➌ Load file eng to HBase
  - ➍ This table named {STUDENT\_ID}\_eng with column-family **grade** and qualifier **name,eng**

# Problem Description

- Part 2 - Query data by Hive or Pig from HBase

Write a script to satisfy the following statement

- 1 Load table from HBase to Hive(Pig)
- 2 JOIN these two tables to a new table named {STUDENT\_ID}\_score
- 3 CREATE a new column named **avg** ( $= (\text{eng} + \text{math})/2$  )
- 4 Find #students failed (  $\text{avg} < 60$  )
- 5 Show the name of top 5 students



# Outline I

- 1 Overview
- 2 HBase
- 3 Hive
- 4 Pig
- 5 Lab Exercise
- 6 Programming Guide**
- 7 Reference

# Login to server

- **Host:** 140.114.91.199 (with 1 master, 8 slaves)
- **Account:** {YOUR\_STUDENT\_ID}
- **Password:** cloud5566 (default)

It's recommended to use ***passwd*** to change your password.

# Grading

- Part 1 - Load data to HBase by API
  - 20% Load data to table {STUDENT\_ID}\_math
  - 20% Load data to table {STUDENT\_ID}\_eng
- Part 2 - Query data by Hive or Pig from HBase
  - 30% Count #students failed
  - 30% Find top 5 students

# Outline I

- 1 Overview
- 2 HBase
- 3 Hive
- 4 Pig
- 5 Lab Exercise
- 6 Programming Guide
- 7 Reference**

# Reference

- <http://hortonworks.com/blog/apache-hadoop-2-is-ga/>
- <http://contest.trendmicro.com/2014/cn/material/hbase.pdf>
- <http://www.datascience-labs.com/hbase/>
- <http://gethue.com/hadoop-tutorial-use-pig-and-hive-with-hbase/>